
What does self-supervision add to engineered features for stellar light curves?

Zijun Liu

University of California, San Diego
br1055@ucsd.edu
liuzijun6688@gmail.com

Yue Ma

Christopher A. Theissen

Department of Astronomy and Astrophysics
University of California, San Diego

Abstract

TESS light curves encode a star’s pulsation, rotation, binarity and oscillations, among other physical properties, but the survey produces far more light curves than can be labelled by hand. Astronomers have engineered a few dozen physically motivated features from them, which together remain a strong baseline. Previous work has trained encoders on light curves and judged them by their own downstream scores, but has not asked what they add to the engineered features. We pre-train a convolutional encoder on unlabelled TESS light curves to predict each window’s latent representation from those of its neighbours. We then freeze it, add its features to the 25 engineered features, and read the combination out with one fixed linear probe on ten tasks: five variability classes, four asteroseismic targets and rotation period. Adding the learned features improves the readout on 7 of 10 tasks over the engineered features alone, while an untrained encoder does not. The combination also outperforms two networks trained end-to-end on the same labelled stars on most tasks. An encoder trained without labels therefore gives astronomers new features that complement the ones they already use rather than replace them.

1 Introduction

A star’s light curve carries its pulsation, rotation, eclipses and oscillations, and with them its age, interior and companions. Wide-field surveys now produce light curves far faster than they can be labelled. TESS [Ricker et al., 2015] returns 2-minute or 20-second [Huber et al., 2022] light curves for hundreds of thousands of pre-selected stars, and full-frame images of every star in the field at cadences that change between mission cycles (30 minutes, then 10, now 200 seconds). LSST [Ivezić et al., 2019] will add *billions* more light curves, albeit at a cadence of days. Self-supervised learning is the natural response: pre-train one encoder on the unlabelled archive, then attach a cheap readout for each task. Encoders of this shape exist for ground-based, Gaia and Kepler light curves [Donoso-Oliva et al., 2025, Zuo et al., 2025, Rui, 2026] and are evaluated by classification accuracy, sometimes against engineered features [Rui, 2026]; a label-free TESS representation [Poznanski, 2026] is evaluated by whether repeat observations of one star lie close in its embedding space. None measures what they add to engineered features.

In this domain, the baseline that a learned encoder must improve on is already strong. Decades of variable-star research have distilled light curves into compact engineered summaries: periodogram peak structure, band-limited power, spectral entropy and robust scatter statistics. Pipelines built on them remain highly competitive. The benchmark of Li et al. [2025] makes the point sharply: across seven expert-labelled variable-star classes, no time-series foundation model they tested strictly surpassed the long-established domain feature baseline on classification. A learned encoder for this data therefore has to answer one question: do its features add predictive signal to the engineered features under the same readout?

We measure that directly. We pre-train an encoder on TESS light curves, freeze it, and add its features to the 25 engineered features under one fixed linear readout. We ask two things: what the learned features add, and whether the combination, which trains only a linear model on the labels, outperforms two networks trained end-to-end on the same labels. Our contributions are:

- **The learned features add to the engineered features rather than replacing them.** Adding the learned features to the engineered ones improves a fixed linear readout beyond 2 SE on 7 of 10 tasks; an untrained encoder gives no such gain, so the gain comes from pre-training, not from merely adding 128 columns.
- **The combined pipeline, read out by a fixed linear model, outperforms the two supervised networks we train end-to-end on the same labels on most tasks.** Against a supervised Conv1D with the encoder’s architecture it is ahead on 6 of 10 tasks and indistinguishable on the other 4; against an MLP on raw flux it is ahead on 9 of 10.

2 Data and method

Data. We use 2-minute SPOC PDCSAP light curves of stars brighter than TESS magnitude $T = 10$ from Sectors 27–101. Pre-training and five of the ten probes (pulsating, eclipsing binary, rotation, transit and rotation period) use a label-enriched subset of 13,470 stars, split 70/15/15 by star: every star labelled as a transit host, eclipsing binary or pulsator that has a usable window, plus a random sample of 10,000 quiet stars, listed in none of our variability catalogues. The pre-training loss never sees a label, but this pool is label-enriched, not a random draw from the survey. The other five probes draw their labels from catalogues that barely overlap this subset and are scored on a second pool of 22,860 stars built the same way (every catalogue positive with a usable window plus 10,000 quiet stars, split 70/15/15). No test star of the first pool enters pre-training. The first pool admits positives only through the transit, eclipsing-binary and pulsating catalogues and screens its quiet stars against the rotation catalogue as well, so every rotator in it carries a second label and it holds no rotation-only star; Appendix J rescores the two rotation tasks on a third pool of 106,284 stars, the corpus outside the first pool, where rotators are a natural 13%. About 8% of the second pool’s test stars (265 of 3,429) were in the pre-training subset; removing them leaves all seven gains in place (Appendix I). Each light curve is split into contiguous segments wherever the time axis breaks, at a missing cadence or a jump longer than five median cadences; segments are median-normalised by their median absolute deviation and sliced into non-overlapping windows of 256 cadences ($256 \times 120 \text{ s} \approx 8.5 \text{ h}$); restricting to one cadence keeps every window the same span of time. Windows containing a missing value are discarded, never interpolated or padded, and segments are never stitched across sector or gap boundaries. Pre-training consumes sequences of exactly 16 consecutive windows ($\approx 5.69 \text{ d}$), cut at a random offset within a segment, and uses no labels.

Encoder. A 1D convolutional variational encoder [Kingma and Welling, 2014] maps each window to a 128-dimensional Gaussian posterior, and a GRU [Cho et al., 2014] predicts the next window’s latent from its predecessors, run both forward and backward over the sequence. The objective is

$$\mathcal{L} = \|\hat{x} - x\|^2 + \beta \text{KL} + \lambda [\|\hat{z}_{t+1} - \bar{z}_{t+1}\|^2 + \|\hat{z}_{t-1} - \bar{z}_{t-1}\|^2] \\ + w [\|\log P_H(\hat{x}) - \log P_H(x)\|^2 + \|\Delta\hat{x} - \Delta x\|^2],$$

where every term is a mean squared error, $\bar{z}$ the stop-gradient target of the forward and backward predictors, P_H the Hann-tapered power spectrum, Δ the first difference, and $(\beta, \lambda, w) = (0.1, 60, 0.3)$, β warmed up linearly over ten epochs. The encoder is trained with six seeds and then frozen; no gradient from any label reaches it.

Representations and readout. A star’s learned features are μ , the mean, over every window of its first segment (16 for 95% of stars, 20 for the rest), of the per-window posterior means. The engineered baseline is a fixed 25-feature vector modelled on the T’DA variability pipeline [Audenaert et al., 2021], computed on the same first segment, 19 from its periodogram and 6 time-domain statistics; Appendix H defines each feature. Readouts are fixed in advance and never tuned on a task, with inputs standardised using training-split statistics alone: class-balanced L_2 logistic regression at $C = 1$ for the detection tasks, ridge for the two regressions with its penalty set by leave-one-out cross-validation. We compare four arms, each the same readout on a different feature set: `features`, μ , `features` $\oplus$ μ , and, as a control, the same fusion built on an *untrained* encoder; the two supervised baselines below are the fifth and sixth arms. Appendix D repeats the first three under XGBoost.

Table 1: Ten-task scorecard. Frozen arms: a fixed logistic or ridge readout, or XGBoost on the same inputs (Appendix D). Bold: best point estimate of the three arms under each frozen readout, with no significance test attached; bold in the appendix tables instead marks a difference resolved beyond its error bar, so the two conventions are not interchangeable. n : test stars; prev.: positive-class prevalence.

task	metric	n (prev.)	frozen linear readout			frozen XGBoost readout			end-to-end supervised	
			features	μ	features $\oplus \mu$	features	μ	features $\oplus \mu$	Conv1D	MLP
pulsating	PR-AUC	2021 (0.107)	0.789	0.806	0.846	0.845	0.821	0.834	0.816	0.748
eclipsing binary	PR-AUC	2021 (0.097)	0.742	0.771	0.775	0.798	0.789	0.807	0.774	0.719
rotation	PR-AUC	2021 (0.089)	0.540	0.559	0.555	0.596	0.544	0.576	0.565	0.526
transit	PR-AUC	2021 (0.060)	0.190	0.144	0.218	0.217	0.223	0.228	0.134	0.109
oscillating giant	PR-AUC	3429 (0.383)	0.920	0.854	0.918	0.924	0.899	0.923	0.875	0.883
solar-like osc.	PR-AUC	3429 (0.122)	0.320	0.336	0.390	0.470	0.409	0.463	0.396	0.275
RGB vs. HeB	ROC-AUC	161 (0.702)	0.758	0.660	0.708	0.746	0.725	0.763	0.530	0.762
$\nu_{\max}$	R^2	1313	0.831	0.802	0.867	0.916	0.831	0.919	0.832	0.834
rotation period	R^2	150	0.703	0.677	0.717	0.695	0.655	0.730	0.597	0.698
flare activity	PR-AUC	3429 (0.089)	0.474	0.453	0.524	0.562	0.485	0.565	0.533	0.444

Supervised baselines. Two end-to-end baselines see exactly the same input windows, star sets and splits. **C1** reuses the encoder’s convolutional trunk unchanged, adds a 128-unit linear layer and a scalar output, and is trained per task with a class-weighted loss on a star-level score, the mean of the star’s window logits computed inside the network so that the gradient flows through the mean. **C2** replaces the convolutional trunk with a dense one and changes nothing else, so C1 versus C2 isolates the convolutional inductive bias. Both use one pre-registered regularisation setting for all tasks, select on a validation split no probe has ever touched, and are trained with six seeds.

Tasks and metrics. We score ten tasks with one probe per physical quantity: pulsating variables [Gao et al., 2025], eclipsing binaries [Prša et al., 2022], rotation and rotation period [Boyle et al., 2026], transits [Guerrero et al., 2021], flare activity [Seli et al., 2025, Vida et al., 2021], oscillating red giants and their $\nu_{\max}$ [Hon et al., 2021], solar-like oscillators [Hatt et al., 2023], and RGB versus helium-burning discrimination [Sreenivas et al., 2026]. Here $\nu_{\max}$ is the frequency of maximum oscillation power. Detection tasks are scored by the area under the precision–recall curve (PR-AUC), whose floor is the positive-class prevalence rather than 0.5; prevalence is therefore quoted with every detection score. The two regressions are scored by R^2 .

Uncertainties. All error bars are $2\times$ the standard error of the mean over seeds. For the μ arms the seeds are independent pre-training runs (six); for C1 and C2 they are initialisation and shuffling seeds (six); the engineered arm is deterministic. The seeds of the μ arms and of the supervised arms are unrelated, so differences between them are treated as unpaired and their bands combined as $\sqrt{\text{SE}_a^2 + \text{SE}_b^2}$. A difference counts only when it exceeds the combined 2 SE band. These bands measure re-training spread only; a paired bootstrap over the test stars (1,000 resamples, seeds averaged; Appendix G) checks that no verdict hinges on the particular test stars.

3 Results

The learned features add to the engineered features. Table 1 and Figure 1a give the comparison. Adding μ to the 25 features improves the fixed readout beyond 2 SE on 7 of 10 tasks, led by solar-like oscillators $+0.071$, pulsating $+0.057$ and flare activity $+0.051$, and reaching $+0.013$ (R^2) on rotation period. Rotation is positive but just short of the bar ($+0.016$ against 0.016). On the rotation pool of Appendix J, where the two classes differ in rotation and nothing else, the rotation gain is $+0.055 \pm 0.006$ but an untrained encoder also gains $+0.029$, so about half of it is what any 128 extra columns buy on 74,000 training stars; the rotation-period gain there is $+0.034 \pm 0.002$ with the untrained control at zero. Neither verdict moves the count. The two negative tasks are the ones where 128 extra columns cost more than they add: the 25 features already cover the task (oscillating giants, -0.003 on a baseline of 0.920) or the training set is too small to fit 128 more weights (RGB versus helium-burning, -0.050 on 755 training stars). On both, an untrained encoder loses by about the same amount as the trained fusion, within the error band; that is what redundant columns do to a linear fit, not a sign that μ carries misleading content (Appendix C). Across all ten tasks the untrained

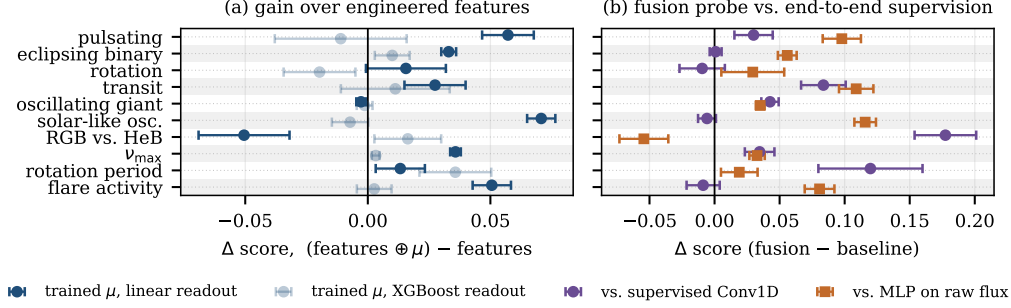


Figure 1: The two claims. **(a)** What the learned features add to the engineered features, under the fixed linear readout of Table 1 (solid) and under an XGBoost readout on the same inputs (faded; Appendix D). **(b)** The fusion probe against the two supervised baselines. Bars are 2 SE.

control is negative on seven and never exceeds $+0.016$ (Table 2), so the gains are not what any 128 extra columns would buy on this feature set; they are specific to what pre-training learned. μ on its own does not replace the features (ahead on only 4 of 10 tasks alone), as Li et al. [2025] also find for time-series foundation models; we therefore report fusion rather than substitution [Makinen et al., 2024].

The fusion probe outperforms the two supervised networks on most tasks. Against C1, which shares the convolutional trunk, labels and input windows but is trained end-to-end, the fusion probe is ahead on 6 of 10 tasks and indistinguishable from it on the other 4 (Figure 1b). Against C2, an MLP on raw flux, it is ahead on 9 of 10 and behind on one (RGB versus helium-burning, 161 stars). C1 is not weak: it is ahead of the engineered arm on 5 of 10. A paired star-bootstrap (Appendix G) excludes zero on 5 of the 7 fusion gains, 5 of the 6 leads over C1 and 7 of the 9 over C2. This holds for this architecture on 669 to 16,002 labelled stars per task, not for supervision in general.

The GRU prediction term raises the gain on every task. Removing the GRU prediction term, all else fixed and the same six seeds, reduces the fusion gain on all 10 tasks, and by more than 2 SE on 6 when the two arms are differenced seed by seed (Table 2). Without it the gain falls by more than half on 5 of the 8 tasks where fusion gains.

4 Discussion

Pick the encoder by a held-out probe, not by its loss. Across the ten pre-training recipes of one sweep, a lower validation reconstruction loss went with a *lower* probe score: the Spearman ρ between negated loss and PR-AUC runs from -0.58 to -0.84 over four tasks (Figure 2). The recipe with the lowest loss is the one without the prediction term on all four, and choosing it would cost 0.021 – 0.049 PR-AUC. On the family our encoder belongs to the anti-correlation does not survive as such — it holds on two of the four tasks and reverses on two — but the loss optimum is still never the probe optimum (Appendix E). The same objective carries a second cost: it develops a localised artifact inside each window, and none of the interventions we tried on four axes cleared both pre-registered acceptance criteria, so the recipe we ship does so by the stated fallback rule rather than by winning (Appendix F). An encoder here therefore has to be chosen by a held-out probe, not by its own loss.

Limitations and next steps. *Readout.* The four linear-readout arms (features, μ , features $\oplus \mu$, and the untrained control) share one linear model. It has few parameters, so it cannot rescue weak features, and it is identical across the four, so the comparison is between feature sets. Our claim is confined to these 25 features and a linear readout: under XGBoost on the same inputs the gain exceeds 2 SE on 4 of 10 tasks and is negative on one; where the nonlinear readout raises the features-only score, the μ gain falls, with a rank correlation of 0.92 across the ten tasks (Appendix D). *Data.* Every arm, the supervised baselines included, sees one 5.7-day segment of a bright 2-minute star. The fainter, far more numerous stars in the full-frame images, and light curves spanning several sectors,

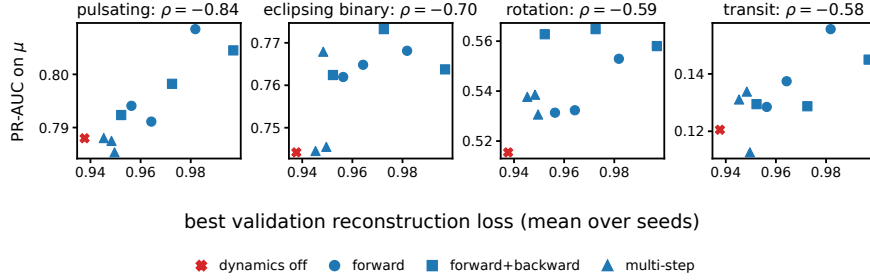


Figure 2: Validation reconstruction loss against downstream probe score, one point per pre-training recipe and one panel per task. The ten recipes share corpus, architecture and loss terms and differ only in the GRU prediction term (off, forward, forward+backward, multi-step, the last three at three weights), four seeds each. x : best post-warm-up validation reconstruction loss, averaged over the recipe’s seeds; y : PR-AUC of the fixed linear probe on μ ; ρ : Spearman’s rank correlation between negated loss and score. The loss-optimal recipe is the dynamics-free one (red) on all four tasks, and it is the worst or next-to-worst probe on each. Appendix E repeats the measurement on the family this paper’s encoder belongs to.

are the obvious next experiments. *Labels*. Our negatives are stars absent from the label catalogues, not stars verified to lack the signal, so catalogue incompleteness adds positive-unlabelled label noise and the detection scores may be conservative. The flare label marks stars that have flared somewhere in their TESS data, not a flare inside the 5.7-day segment we encode, so the task is recognising a flaring star, not finding a flare. *Label scarcity*. When the readout is trained on a few percent of the labels, the advantage narrows on 7 of 10 tasks, so how much pre-training helps when labels are scarce remains open. *Discovery*. Every task here has a catalogue; whether μ also separates kinds of variability that no catalogue labels yet, for instance as clusters in the latent space, is untested. We leave it as an exploratory direction rather than a result: with no labels there is nothing to score such a separation against, so it needs a design of its own.

Reproducibility and disclosure. §2 and Appendix H give the data, objective, features and readouts; the encoder has four convolutional stages (widths 32/64/128/256, kernel 5, each halving length), a 128-d latent and a one-layer GRU of width 256. Code and score tables will be released on publication; generative AI assisted with analysis code and prose editing.

References

- Jeroen Audenaert, James S. Kuszlewicz, Rasmus Handberg, et al. TESS data for asteroseismology (T²DA) stellar variability classification pipeline: Set-up and application to the Kepler Q9 data. *The Astronomical Journal*, 162(5):209, 2021. arXiv:2107.06301.
- Timothy R. Bedding, Benoit Mosser, Daniel Huber, et al. Gravity modes as a way to distinguish between hydrogen- and helium-burning red giant stars. *Nature*, 471:608, 2011. arXiv:1103.5805.
- Andrew W. Boyle, Luke G. Bouma, and Andrew W. Mann. The TESS all-sky rotation survey: Periods for 1,046,317 stars within 500 pc. *arXiv e-prints*, 2026. arXiv:2603.05586.
- Kyunghyun Cho, Bart van Merriënboer, Caglar Gulcehre, et al. Learning phrase representations using RNN encoder–decoder for statistical machine translation. In *Conference on Empirical Methods in Natural Language Processing*, 2014. arXiv:1406.1078.
- Cristobal Donoso-Oliva, Ignacio Becker, Pavlos Protopapas, et al. Astromer 2. *arXiv e-prints*, 2025. arXiv:2502.02717.
- Xinyi Gao, Xiaodian Chen, Shu Wang, and Jifeng Liu. Classification of periodic variable stars from TESS. *The Astrophysical Journal Supplement Series*, 276(2):57, 2025. arXiv:2412.06175.

- Natalia M. Guerrero, Sara Seager, Chelsea X. Huang, et al. The TESS objects of interest catalog from the TESS prime mission. *The Astrophysical Journal Supplement Series*, 254(2):39, 2021. arXiv:2103.12538.
- Emily Hatt, Martin B. Nielsen, William J. Chaplin, et al. A catalogue of solar-like oscillators observed by TESS in 120-second and 20-second cadence. *Astronomy & Astrophysics*, 669:A67, 2023. arXiv:2210.09109.
- Marc Hon, Daniel Huber, James S. Kuszlewicz, et al. A ‘quick look’ at all-sky galactic archeology with TESS: 158,000 oscillating red giants from the MIT quick-look pipeline. *The Astrophysical Journal*, 919(2):131, 2021. arXiv:2108.01241.
- Daniel Huber, Timothy R. White, Travis S. Metcalfe, et al. A 20-second cadence view of solar-type stars and their planets with TESS: Asteroseismology of solar analogs and a re-characterization of π Men c. *The Astronomical Journal*, 163(2):79, 2022. arXiv:2108.09109.
- Željko Ivezić, Steven M. Kahn, J. Anthony Tyson, et al. LSST: From science drivers to reference design and anticipated data products. *The Astrophysical Journal*, 873(2):111, 2019. arXiv:0805.2366.
- Diederik P. Kingma and Max Welling. Auto-encoding variational bayes. In *International Conference on Learning Representations*, 2014. arXiv:1312.6114.
- Weijian Li, Hong-Yu Chen, Nabeel Rehemtulla, et al. StarEmbed: Benchmarking time series foundation models on astronomical observations of variable stars. *arXiv e-prints*, 2025. arXiv:2510.06200.
- T. Lucas Makinen, Ce Sui, Benjamin D. Wandelt, et al. Hybrid summary statistics. *arXiv e-prints*, 2024. arXiv:2410.07548.
- Dovi Poznanski. A useful representation of TESS light curves. *arXiv e-prints*, 2026. arXiv:2605.05324.
- Andrej Prša, Angela Kochoska, Kyle E. Conroy, et al. TESS eclipsing binary stars. I. short-cadence observations of 4584 eclipsing binaries in sectors 1–26. *The Astrophysical Journal Supplement Series*, 258(1):16, 2022. arXiv:2110.13382.
- George R. Ricker, Joshua N. Winn, Roland Vanderspek, et al. The transiting exoplanet survey satellite. *Journal of Astronomical Telescopes, Instruments, and Systems*, 1(1):014003, 2015. arXiv:1406.0151.
- Yicheng Rui. Domain-informed multi-view self-distillation for astronomical light-curve representation learning with JEPA. *arXiv e-prints*, 2026. arXiv:2606.28446.
- Bálint Seli, Krisztián Vida, Katalin Oláh, et al. Stellar flare morphology with TESS across the main sequence. *Astronomy & Astrophysics*, 694:A161, 2025. arXiv:2412.12989.
- K. R. Sreenivas et al. Global asteroseismology of 19,000 red giants in the TESS continuous viewing zones. *arXiv e-prints*, 2026. arXiv:2604.00498.
- Krisztián Vida, Attila Bódi, Tamás Szklénár, and Bálint Seli. Finding flares in Kepler and TESS data with recurrent deep neural networks. *Astronomy & Astrophysics*, 652:A107, 2021. arXiv:2105.11485.
- Xiaoxiong Zuo, Yihan Tao, Yang Huang, et al. FALCO: A foundation model of astronomical light curves for time-domain astronomy. *arXiv e-prints*, 2025. arXiv:2504.20290.

A Pre-training and readout details

What one pre-training example is. The packed pre-training corpus is the 70% training split of the first pool: 9,428 stars, 35,854 segments and 582,352 windows of 256 cadences. The encoder is monitored on the pool’s 15% validation split (2,021 stars, 121,328 windows), which no probe ever sees. One example is one segment. Each epoch it yields a fresh run of 16 consecutive windows beginning at a uniformly random offset inside that segment, so a long segment shows a different 5.69-day stretch every epoch; a segment with fewer than 16 windows never entered the pack. Validation uses offset 0 and every segment, so the monitored quantity does not move with the sampling. A batch is 256 sequences, giving 140 optimiser steps per epoch.

The first-difference term. In the objective of §2, $\Delta x_t = x_{t+1} - x_t$ is the first difference along a window, and $\|\Delta \hat{x} - \Delta x\|^2$ is the mean squared error between the reconstruction’s cadence-to-cadence changes and the input’s. It is there because the point-wise term $\|\hat{x} - x\|^2$ is already small for a smooth curve drawn through a noisy one: the fastest components of a light curve carry little amplitude, so a decoder that low-passes its input pays almost nothing for discarding them. Differencing is a high-pass operation, so the same error re-enters the objective weighted towards high frequency and the smoothed reconstruction is penalised. P_H plays the same role in the frequency domain; the two share the weight w and are summed with equal weight inside it.

Optimisation. Adam at learning rate 3×10^{-4} , cosine-annealed to 3×10^{-6} over the 100-epoch cap, gradient-norm clipping at 1.0, mixed precision. β rises linearly from 0 to its target 0.1 over the first ten epochs and stays there; the shipped recipe uses no free-bits floor. Training stops early once ten consecutive epochs improve no tracked checkpoint, a counter that is armed only after the warm-up.

Which checkpoint is frozen. Only post-warm-up epochs are eligible. During the warm-up β and the KL term are still moving, and the scheduled- β total is minimised at $\beta = 0$, that is by an untrained model, so an unrestricted argmin can select the warm-up transient rather than a fit. Among the eligible epochs we freeze the one minimising recon + $w \cdot$ aux on the validation split. The prediction term is deliberately absent from that metric: it is the quantity a dynamics ablation changes, so including it would select each arm of that comparison by a different rule.

Readouts. Logistic regression is scikit-learn’s LogisticRegression with L_2 penalty, $C = 1$, class_weight="balanced" and at most 2,000 iterations. Ridge is RidgeCV over ten log-spaced penalties from 10^{-2} to 10^3 , chosen by leave-one-out cross-validation. Every input column is z-scored with the training split’s own mean and standard deviation, and the same two numbers are applied unchanged to validation and test. For the fusion arm the 25 features and the 128 latent means are column-concatenated first and the 153-column matrix is standardised as a whole, so no arm gets a different scaling of the columns it shares with another.

The supervised baselines. C1 and C2 are trained with AdamW at learning rate 3×10^{-4} and weight decay 10^{-4} , dropout 0.2 before the task head only (the trunk is unregularised, as the encoder’s is), gradient clipping at 1.0, mixed precision, batches of 64 stars ($\approx 1,040$ windows, since the loss is per star), at most 60 epochs and patience 10. The positive class enters the loss with weight $n_{\text{neg}}/n_{\text{pos}}$ computed on the training split. One regularisation setting is fixed in advance for all ten tasks and no per-task tuning is done. Each run is selected on the best validation score in that task’s own headline metric rather than on validation loss: at the prevalences here the validation-loss minimum and the validation-metric maximum routinely fall many epochs apart, and selecting on the loss would depress the baseline we are trying to beat.

B The two star pools, and what a quiet star is

Each task is scored on the pool its catalogue populates, and every arm on a task sees the identical star list, so the two definitions below fix the population but never a comparison.

The first pool (13,470 stars). Every star in the corpus with at least one usable window that carries a transit, eclipsing-binary or pulsating label: 3,470 stars, assigned to a single split stratum rarest class first (813 transit, 1,226 eclipsing binary, 1,431 pulsating). To them are added 10,000 *quiet* stars drawn uniformly at random, without replacement, from the corpus stars that appear in none of our five v1 catalogues – transit, eclipsing binary, pulsating, rotation and flare. The whole pool is split 70/15/15 by star within each stratum, and the encoder is pre-trained and monitored on this pool’s training and validation splits.

Note what the two clauses do differently. A star becomes a positive member only through the three detection catalogues; rotation and flare appear only as exclusions on the quiet draw. A rotator that is not also a transit host, eclipsing binary or pulsator is therefore neither admitted as a positive nor drawn as a quiet star, and is absent from this pool altogether. Appendix J measures what this costs the two rotation rows.

The second pool (22,860 stars). The four external catalogues behind the asteroseismic and rotation-period tasks are joined to the corpus into one table of 11,007 stars, each row a star that at least one of them covers. Those 11,007, plus the 1,853 flare-labelled corpus stars not already among them, are the pool’s members (12,860). To them are added a fresh draw of 10,000 quiet stars: corpus stars that carry no transit, eclipsing-binary, pulsating or rotation label, have no flare label, and are absent from the external table. The split is again 70/15/15 by star, with the stratum taken rarest first over RGB-versus-helium-burning, Kounkel rotation period, flare, solar-like oscillator, oscillating giant and quiet.

The two screens are not symmetric. The first pool’s quiet draw is screened against the five v1 catalogues only; it never consults the four external ones, so some of its quiet stars are oscillating red giants or solar-like oscillators. The second pool’s quiet draw is screened against both sets. Nothing printed here compares a star in one pool with a star in the other, so the asymmetry costs no verdict; what it does mean is that a negative in either pool is a star *absent from the relevant catalogues*, not a star examined and found to lack the signal. The external table is itself a union of four detection lists, so a star it covers with a zero for oscillating giants is a star Hatt, Sreenivas or Kounkel listed and Hon et al. [2021] did not, which is the same kind of negative. This is the label limitation §4 states, given exactly.

C The two tasks where fusion does not help

task	features	fusion Δ	2 SE	untrained Δ	n_{test}
oscillating giants	0.920	−0.003	0.002	−0.003	3,429
RGB vs. He-burning	0.758	−0.050	0.019	−0.038	161

Oscillating giants: the 25 features already cover the task. The labels of Hon et al. [2021] come from a convolutional classifier reading a 128×128 image of each star’s power spectrum, which fires on a visible oscillation power excess at $6 \lesssim \nu_{\text{max}} \lesssim 250 \mu\text{Hz}$. Our engineered vector encodes close to that same quantity: eight log-spaced band powers spanning 0.3–360 c/d, plus the spectral centroid and entropy, locate the excess and describe its concentration. The features accordingly reach 0.920, and the trained and untrained fusion arms then lose the same 0.003: 128 collinear columns diluting a linear fit, with no trained-versus-untrained difference left to attribute to μ .

RGB versus helium-burning: no arm resolves the discriminant. Sreenivas et al. [2026] assign the evolutionary state with a convolutional classifier reading the background-corrected power spectrum *folded at $\Delta\nu$* ; the physics underneath is the dipole mixed-mode period spacing $\Delta\Pi_1$ [Bedding et al., 2011]. Neither is accessible at our baseline. The encoder sees one 256-cadence window at a time, so the finest frequency structure available to μ is $\Delta f = 32.6 \mu\text{Hz}$, while $\Delta\nu$ for these stars is $\approx 1.5\text{--}18 \mu\text{Hz}$; the comb itself is unresolved, let alone a fold of it. The engineered features, computed on the whole 5.69 d segment, reach $2.0 \mu\text{Hz}$, still coarser than the mixed-mode spacing $\Delta\Pi_1\nu^2 \approx 0.2\text{--}0.6 \mu\text{Hz}$ near $\nu_{\text{max}} = 50 \mu\text{Hz}$. No arm on this row measures the quantity the label encodes; the 0.758 the features reach rides on correlates of evolutionary state rather than on its discriminant. The row behaves accordingly: it is the only one where C2 (0.762) is ahead of the fusion probe while C1 reaches only 0.530 on 755 training stars, and it carries the smallest test set in the study.

Why 128 redundant columns cost a linear fit. A linear readout spends its training stars estimating one coefficient per column, and each estimate carries its own error into the prediction whether or not the column carries signal. On the RGB-versus-helium-burning row the features arm fits 25 coefficients on 755 training stars and the fusion arm 153, so the stars-per-coefficient ratio falls from about 30 to about 5 while, by the resolution argument above, the added columns supply nothing the label depends on. The variance a column adds scales with how few stars are available to pin it down; the signal it adds does not. A column is therefore worth its place only when what it adds exceeds that cost, and on this row and on oscillating giants it does not. The logistic readout’s L_2 penalty is fixed at $C = 1$ for every task rather than tuned per task, so it damps this effect but does not remove it. The prediction that follows is the one the table above shows: an untrained encoder’s 128 columns should

cost about the same as a trained one's, since the argument never mentions what the columns contain, and they do (-0.038 against -0.050 , and -0.003 against -0.003).

D The gain under a nonlinear readout

Table 2: What 128 added columns change, relative to the 25 features alone, under each frozen readout (mean over six encoder seeds; absolute scores are in Table 1). $+\mu$: our encoder. No prediction term: the same recipe trained without the GRU prediction term. Retained: the share of the full μ gain that survives that removal, printed only where the full gain is a gain, and negative where the gain changes sign rather than shrinking. Untrained: a randomly initialised encoder, which is a single encoder and carries no band. Bold: beyond 2 SE — unlike Table 1, where bold marks the best point estimate and implies no test. The last row counts tasks won, tied and lost at that bar.

task	metric	linear readout				XGBoost readout	
		$+\mu$	$+\mu$, no prediction term	retained	$+\text{untrained } \mu$	$+\mu$	$+\text{untrained } \mu$
pulsating	PR-AUC	+0.057	+0.051	0.89	+0.016	-0.011	-0.003
eclipsing binary	PR-AUC	+0.033	+0.012	0.38	-0.010	+0.010	-0.003
rotation	PR-AUC	+0.016	-0.011	-0.69	-0.026	-0.020	-0.034
transit	PR-AUC	+0.027	+0.018	0.66	+0.004	+0.011	-0.027
oscillating giant	PR-AUC	-0.003	-0.006	-	-0.003	-0.001	-0.008
solar-like osc.	PR-AUC	+0.071	+0.024	0.33	-0.007	-0.007	-0.026
RGB vs. HeB	ROC-AUC	-0.050	-0.052	-	-0.038	+0.016	-0.031
$\nu_{\max}$	R^2	+0.036	+0.010	0.29	-0.001	+0.003	-0.005
rotation period	R^2	+0.013	-0.005	-0.41	-0.057	+0.036	-0.004
flare activity	PR-AUC	+0.051	+0.033	0.65	+0.006	+0.003	+0.006
win / tie / loss		7 / 1 / 2				4 / 5 / 1	

The retained column is the per-task form of the body's ablation count. On the eight tasks where fusion gains, removing the GRU prediction term leaves between -0.69 and 0.89 of that gain, below half on five of them, and on two (rotation and rotation period) the remainder is negative, so the ablated encoder's columns are worse than not adding them at all. We give the eight fractions rather than one summary number because the ten gains are measured in PR-AUC, ROC-AUC and R^2 and cannot be added together.

The XGBoost readout takes the same inputs as the linear one. For each task and seed it is selected over a grid of maximum depth $\{3, 5\}$ and learning rate $\{0.05, 0.1\}$, with up to 1,000 rounds and early stopping after 50 rounds without improvement on a held-out 25% of the training stars; the test stars play no part in the selection. Under XGBoost the features-only arm also varies with the seed, through the validation split, so the gain is differenced seed by seed.

Under XGBoost the μ gain (features $\oplus \mu$ minus features alone) is smaller than under the linear readout on most tasks. XGBoost also scores higher than the linear readout on the 25 features alone on eight of ten tasks, and the two changes go together. On solar-like oscillators the features-only score rises from 0.320 to 0.470 and the μ gain falls from $+0.071$ to -0.007 . On RGB versus helium-burning, where XGBoost does not raise the features-only score, the μ gain is larger under XGBoost. Across the ten tasks, the rise in the features-only score and the fall in the μ gain have a Spearman rank correlation of 0.92. This is consistent with part of the information μ supplies to a linear readout being already present in the 25 features, in a form only a nonlinear readout can use.

Table 2 also gives the same comparison with an untrained encoder in place of the trained one. Under XGBoost its 128 columns lower the score on nine of ten tasks and the trained μ scores above it on eight, so the gains that survive come from pre-training, not from adding columns. XGBoost on the 25 features alone also scores above the linear fusion probe on 7 of 10 tasks (Table 1). The body's claim is about what μ adds under a fixed linear readout, not that the linear fusion probe is the strongest classifier available on these features.

E Validation reconstruction against downstream score

Figure 2 in §4 is half of this measurement; Figure 3 is the other half, and the two differ in which family of recipes is varied.

In Figure 2 the loss-optimal recipe is the dynamics-free one on all four tasks, and it is the worst or next-to-worst probe on each, so selecting on the loss costs 0.021–0.049 PR-AUC against the probe-optimal recipe. The trend is not carried by that one point: with it removed, ρ still runs from -0.43 to -0.87 . Those recipes are an earlier sweep than our encoder – they carry an earlier spectral auxiliary term – so Figure 3 repeats the estimator on the family we ship: 19 recipes built on the final encoder that vary the form, weight and floor of the auxiliary term, six seeds each (two more recipes in which one seed collapsed are excluded; we never drop a seed).

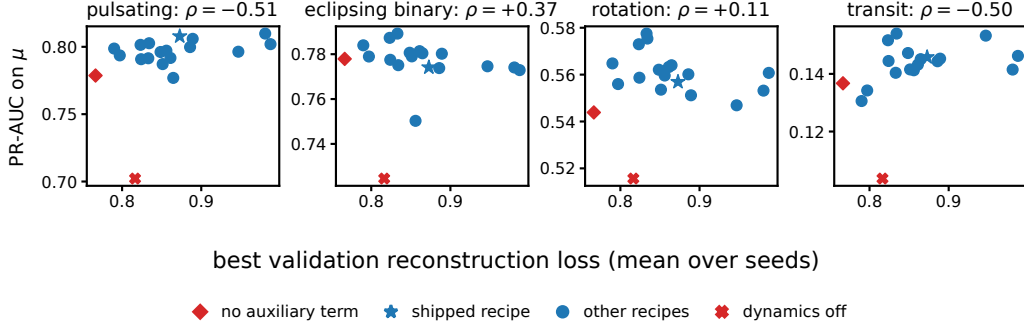


Figure 3: The same estimator as Figure 2, on the family our encoder belongs to: 19 recipes on the final encoder that vary the form, weight and floor of the spectral auxiliary term, six seeds each. The diamond has no auxiliary term, the cross is the one dynamics-off recipe, the star is the recipe we use. x : best post-warm-up validation reconstruction loss; y : PR-AUC of the fixed linear probe on μ ; ρ : Spearman’s rank correlation between negated loss and score.

There the anti-correlation does not persist as such: ρ is -0.51 on pulsating and -0.50 on transit, $+0.37$ on eclipsing binaries and $+0.11$ on rotation. What persists is that the loss is not a selector: its optimum is again the recipe with the weakest competing term (the diamond, no auxiliary term at all), which is never the probe optimum, and trusting it costs 0.011–0.034 PR-AUC. So the body’s figure shows an anti-correlation on every task; this one shows it surviving on two tasks and reversing on two; in neither does the loss choose. Both are statements across recipes only: within one recipe, across its seeds, no label-free curve metric we tried (reconstruction, reconstruction plus auxiliary, active latent units) ranks the seeds, every such correlation sitting within ± 0.3 of zero.

F The within-window artifact, and the two gates that governed it

Why the objective carries a tapered power spectrum. A point-wise reconstruction error prices a light curve by amplitude, and in these data amplitude sits mostly in the slow component, so a decoder can score well on it while discarding the periodic structure the tasks depend on. The two auxiliary terms put the pressure back where the point-wise term does not reach: the first difference in time, and the log power spectrum across frequency. The spectrum is taken on a Hann-tapered window rather than on the raw one, for the reason the rest of this appendix gives.

What the decoder found instead. Every term in the objective is computed inside a single 256-cadence window; nothing spans a boundary. A narrow impulse raises log-power at every frequency at once, so a decoder can lower the log-spectrum error by planting one spike rather than by reproducing the star’s spectrum. Where the cheapest spike sits is decided by the taper, a fixed function of position within the window: with a rectangular window it is the window edge, and under a Hann taper it is the window centre, where the taper weight is largest. Figure 4 shows both, in flux and in the position-resolved error. Laid end to end, the spikes form one train with a period of exactly one window, which is a property of the display and not of the objective.

Measured as the per-position reconstruction error divided by its own interior median (positions 16–240), maximised over position within each seed and then averaged over six seeds, the rectangular recipe reaches $31.7\times$ at the window edge and the tapered recipe $11.8\times$ near the centre. The taper

moved the defect and shrank it; it did not remove it. An earlier claim that it had was based on a measurement taken only at the edge, and is withdrawn.

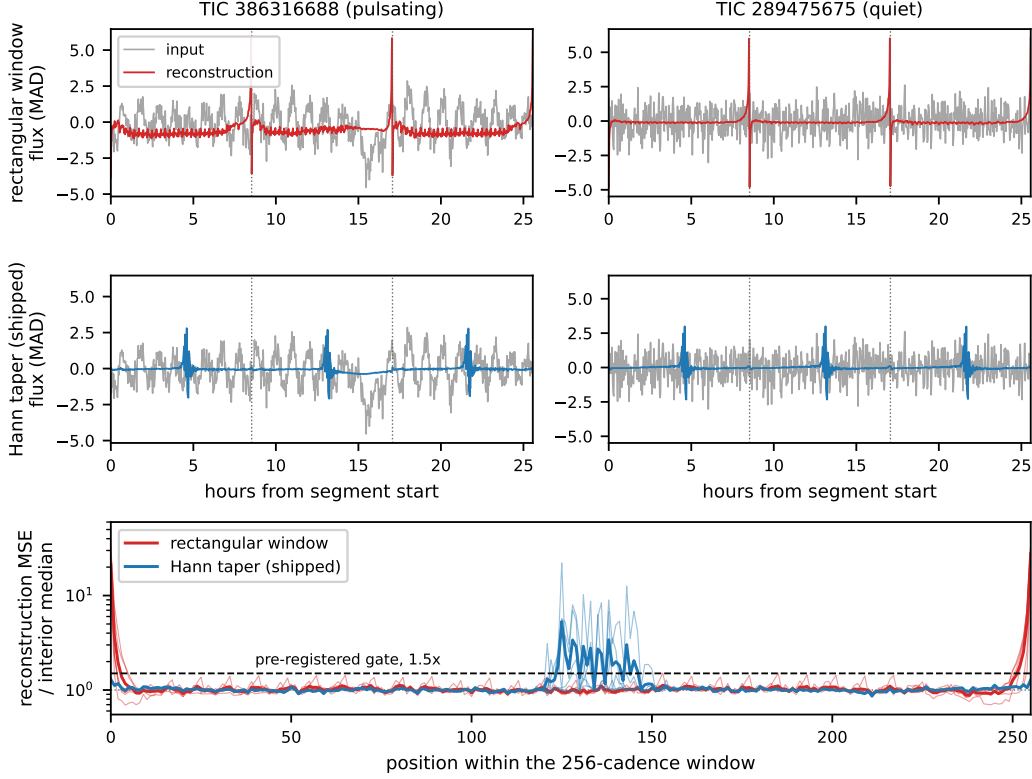


Figure 4: The artifact, in two stars and in the average. **Top two rows:** three consecutive 256-cadence windows of one pulsator and one quiet star, input flux in grey against each recipe’s reconstruction; dotted lines mark the window boundaries, and both rows of a column share one y -scale. The two recipes differ only in the taper used for the auxiliary spectral term. Neither decoder reproduces the cadence-scale scatter, which a 128-dimensional bottleneck cannot carry; what differs is where each one puts a spike it does not need. **Bottom:** reconstruction error by position within the window, divided by that seed’s median over positions 16–240. Thin lines are the six pre-training seeds, thick lines their mean. The dashed line is the pre-registered artifact gate. The seed mean understates the tapered recipe’s peak because its position moves between seeds (125–143), which is why the gate is read as a maximum within each seed.

The two pre-registered acceptance criteria. Both were registered on 2026-08-15, before any of the interventions below were written.

Artifact gate. The within-window impulse ratio – the quantity plotted in the lower panel of Figure 4 – taken as the maximum over *all* positions rather than at the edge alone, at most $1.5\times$ at full auxiliary weight, over six seeds, *and* a visual sign-off on a fresh random sample of reconstructed light curves. The number never passes on its own: reading an edge-only number is how the withdrawn claim above was made in the first place.

No-regression gate. No task among eclipsing binaries, pulsators, rotation and transits worse than the incumbent recipe by more than 2 SE, at both readouts, paired over six seeds. It is the guard against a clean but useless encoder: a recipe that abolishes the impulse and loses the probe is not a fix.

What we tried, and why the tapered recipe is the one we ship. We varied the auxiliary term along four axes: its *mechanism* (remove it entirely; average the spectrum over a family of Slepian tapers so that no single position is cheapest; add a penalty on the excess kurtosis of the residual; both

together; put a floor under the spectral sub-term), its *amount* (six weights of the best mechanism, from 0.0125 to 0.20), the *level* of that floor (rebuilt as a two-sided hinge, which unlike a clamp keeps a restoring gradient, at three measured values), and finally a *floor on the KL term*, which was the remaining untried lever. No cell cleared both gates. The ones that removed the impulse lost the probe; the ones that kept the probe kept a measurable impulse. The nearest miss reached 1.55 ± 0.05 , which is 2.4 SE above the gate; its repaired variant reached 1.47 over six seeds and 1.55 over twelve, and regressed on pulsators at both readouts.

Because no candidate passed, the rule fixed before the ladder ran leaves the incumbent in place. The encoder this paper uses is therefore the Hann-tapered recipe by a pre-registered fallback, not because it won a comparison. Two things make that defensible rather than merely procedural. The artifact was separately measured not to cost the downstream probe – ablating it moved $\nu_{\max}$ by $+0.0001 \pm 0.0037$ – and closing the channel, which was also a test of whether the validation loss becomes a usable selector once it can no longer be lowered by the trick, did not make it one (Appendix E).

G Paired bootstrap over test stars

The 2 SE bands in the body come from re-training and say nothing about whether the particular test stars favoured one arm. For every printed contrast we therefore resample each task’s test set with replacement 1,000 times, rescore every arm on the identical resampled list from its saved per-star predictions (no model is retrained), average each arm over its seeds, and take the difference. Table 3 gives the 95% percentile interval of that difference and the share of resamples in which the fusion arm is ahead; intervals that exclude zero are bold. The check tightens no verdict and reverses none: every row the body calls ahead has the fusion arm ahead in at least 65% of resamples. It does loosen the smallest tasks – rotation period (150 test stars) and RGB versus helium-burning (161) account for four of the ten rows where a resolved 2 SE verdict becomes an interval that touches zero.

Table 3: Paired star-bootstrap of the four contrasts the body counts. Each entry is the 95% interval of (left arm – right arm) over 1,000 resamples of the test stars, seeds averaged, and the share of resamples with the difference above zero ($f_{>0}$; 1.00 means the left arm is ahead in every resample, 0.00 that it is behind in every resample). Bold: interval excludes zero.

task	fusion – features		fusion – C1		fusion – C2		fusion – dynamics-off fusion	
	95% CI	$f_{>0}$	95% CI	$f_{>0}$	95% CI	$f_{>0}$	95% CI	$f_{>0}$
pulsating	[+0.012, +0.105]	0.99	[−0.008, +0.073]	0.94	[+0.045, +0.150]	1.00	[−0.023, +0.036]	0.69
eclipsing binary	[+0.005, +0.062]	0.99	[−0.024, +0.025]	0.52	[+0.022, +0.094]	1.00	[+0.004, +0.038]	0.99
rotation	[−0.026, +0.057]	0.77	[−0.056, +0.040]	0.38	[−0.023, +0.079]	0.87	[−0.005, +0.052]	0.94
transit	[−0.025, +0.065]	0.85	[+0.041, +0.124]	1.00	[+0.046, +0.174]	1.00	[−0.016, +0.035]	0.78
oscillating giants	[−0.010, +0.005]	0.21	[+0.028, +0.057]	1.00	[+0.021, +0.051]	1.00	[−0.002, +0.008]	0.87
solar-like osc.	[+0.044, +0.097]	1.00	[−0.037, +0.022]	0.34	[+0.079, +0.151]	1.00	[+0.025, +0.067]	1.00
RGB vs. HeB	[−0.087, −0.012]	0.00	[+0.099, +0.247]	1.00	[−0.126, +0.020]	0.08	[−0.037, +0.039]	0.52
$\nu_{\max}$	[+0.028, +0.043]	1.00	[+0.024, +0.045]	1.00	[+0.018, +0.047]	1.00	[+0.020, +0.031]	1.00
rotation period	[−0.038, +0.067]	0.65	[+0.054, +0.186]	1.00	[−0.053, +0.100]	0.69	[−0.021, +0.060]	0.83
flare	[+0.019, +0.078]	1.00	[−0.047, +0.029]	0.32	[+0.027, +0.128]	1.00	[−0.000, +0.035]	0.97

H The 25 engineered features

All 25 features are computed from a star’s first segment: its windows concatenated into one flux series x of 4,096 or 5,120 cadences, already median-normalised by its median absolute deviation. The periodogram is the power spectrum of the mean-subtracted series, $P(f) = |\text{FFT}(x - \bar{x})|^2$ with the zero-frequency bin dropped, and $p(f) = P(f) / \sum_f P(f)$ its normalisation to unit sum. Frequencies are in cycles per day. f_1, f_2, f_3 are the three strongest local maxima of P , strongest first.

Table 4: The engineered feature vector, modelled on the T’DA variability pipeline [Audenaert et al., 2021].

#	feature	definition
1–3	$\log_{10} f_1, \log_{10} f_2, \log_{10} f_3$	frequencies of the three strongest periodogram peaks
4–6	p_1, p_2, p_3	normalised power at each peak, $p_k = P(f_k) / \sum_f P(f)$
7–8	$p_2/p_1, p_3/p_1$	peak-power ratios
9	harmonic ratio	$P(2f_1)/P(f_1)$ at the bin nearest $2f_1$; 0 if $2f_1$ exceeds the Nyquist frequency
10–17	band powers	$\log_{10} \sum_{f \in B_j} p(f)$ for eight log-spaced bands B_j spanning 0.3–360 c/d
18	spectral entropy	$-\sum_f p(f) \ln p(f) / \ln N$, N the number of frequency bins
19	spectral centroid	$\log_{10} \sum_f f p(f)$
20–21	skew, kurtosis	sample skewness and excess kurtosis of x
22	point-to-point scatter ratio	$\text{std}(\Delta x) / \text{std}(x)$, Δ the first difference
23	depth	95th minus 5th percentile of x
24	MAD	median absolute deviation of x
25	IQR	inter-quartile range of x

What “modelled on T’DA” means. Audenaert et al. [2021] classify Kepler Q9 light curves into eight variability classes with an ensemble of four supervised learners and a random-forest metaclassifier. We take none of that: not the classifiers, not the classes, not the training set. What we borrow is the shape of the feature vector its tree-based members read, and the borrowing is partial in ways a reader should be able to check.

Kept. A short list of periodogram peaks with their relative strengths, a harmonic measure of the dominant peak, an entropy measure of the signal, and robust time-domain scatter statistics – skewness, kurtosis and the median absolute deviation, all three of which appear in their Table 1. Our point-to-point scatter ratio $\text{std}(\Delta x) / \text{std}(x)$ is, up to a square root, their variability index η_e , the ratio of the mean squared successive difference to the variance.

Changed. They extract up to six frequencies with up to ten harmonics each by iterative prewhitening – fit a frequency and its harmonics, subtract them, repeat until a significance criterion stops the loop. We take the three strongest local maxima of one segment’s periodogram, with no prewhitening and no significance test, and keep one harmonic ratio $P(2f_1)/P(f_1)$ instead of a harmonic series. Their amplitudes and phases become, here, powers normalised by the total and ratios between them: the flux is median-absolute-deviation normalised per segment, so an absolute amplitude is not a quantity we have, and we drop phases and phase differences entirely. Their FliPer, the mean power above each of four frequency cut-offs, becomes our eight log-spaced band powers spanning 0.3–360 c/d plus a spectral centroid. Their entropy features are computed on the time series (differential entropy, multiscale sample entropy); ours is the Shannon entropy of the normalised power spectrum. Because a packed segment is uniformly sampled and gap-free, its Lomb–Scargle periodogram is the FFT periodogram, so we compute the cheaper one.

Dropped. Everything that needs a phase fold at a reliably determined period – the self-organising-map location, the folded point-to-point statistics and the folded range – because 5.69 days covers few cycles for the longer-period classes in our menagerie. Also dropped: the coherency parameter, zero crossings, the range of the cumulative sum, the Shapiro–Wilk statistic, the variance ratio, the significant-harmonic count and the flux ratio, and the power-spectrum image their deep-learning member reads, which is not a feature vector at all.

The result is 25 numbers, 19 from the periodogram and 6 from the time domain, and it is a baseline rather than a reproduction: it is deliberately the cheap, linear-readable part of a pipeline whose published accuracy came from an ensemble we do not run.

I Evaluation restricted to stars never seen in pre-training

The encoder is fitted on the training stars of the first pool and monitored on its validation stars, so the test stars of the five tasks scored on that pool are unseen by construction. The second pool shares some stars with it. Table 5 removes every test star that was in the pre-training training or validation

split and rescores the saved per-star predictions; no model is refitted. The gain exceeds 2 SE on the same 7 of 10 tasks; the largest change is flare activity, +0.051 to +0.041, and the small loss on oscillating giants (−0.003) becomes a tie (−0.001).

Table 5: The fusion gain (features $\oplus \mu$ minus features, linear readout) on all test stars and on the test stars never seen in pre-training. Bold: beyond 2 SE over six encoder seeds.

task	metric	test stars	unseen test stars	gain, all	gain, unseen
pulsating	PR-AUC	2021	2021	+0.057	+0.057
eclipsing binary	PR-AUC	2021	2021	+0.033	+0.033
rotation	PR-AUC	2021	2021	+0.016	+0.016
transit	PR-AUC	2021	2021	+0.027	+0.027
oscillating giant	PR-AUC	3429	3164	−0.003	−0.001
solar-like osc.	PR-AUC	3429	3164	+0.071	+0.069
RGB vs. HeB	ROC-AUC	161	149	−0.050	−0.048
$\nu_{\max}$	R^2	1313	1218	+0.036	+0.034
rotation period	R^2	150	150	+0.013	+0.013
flare activity	PR-AUC	3429	3164	+0.051	+0.041

J The two rotation tasks on a pool that contains rotation-only stars

The defect. As Appendix B states, the first pool admits a star as a positive only through the transit, eclipsing-binary and pulsating catalogues, and its quiet draw excludes every star the rotation catalogue lists. All 1,138 rotators in that pool are therefore also transit hosts, eclipsing binaries or pulsators (607, 383 and 148 by split stratum), and none of its 10,000 quiet stars rotates. A readout can separate the two classes of the *rotation* task by detecting the co-occurring class, and the *rotation period* task is fitted on the same co-labelled rotators (669 training, 150 test). The ML4PS version of this paper scored both tasks on that pool; the body above still prints those numbers, and this appendix reports the same two tasks on a pool built to remove the defect. The design, the population, the arms, the readouts and the rule for reading the result were fixed in writing before any score on the new pool was computed.

The rotation pool. Every corpus star with at least one usable window that is not in the first pool: 106,284 stars, of which 13,789 carry a rotation period in the TESS All-Sky Rotation Survey [Boyle et al., 2026] and 92,495 do not. By construction the pool contains no transit host, eclipsing binary or pulsator (the first pool took every one), no star that entered pre-training, and no sampling step: rotators are 13.0% of it, against 12.5% of the corpus, so both rotation tasks are scored at the survey’s own prevalence rather than the case-control prevalence of the other pools. It is split 70/15/15 by star, stratified on rotators with period at most five days (8,786), other rotators (5,003) and non-rotators, seed 0. Rotation period keeps the five-day cap of the body. A non-rotator here is a star the survey does not list, not a star examined and found not to rotate; this is the same label limitation that bounds the body’s rotation row. No encoder is run again: the per-star encoder outputs and the 25 engineered features for these stars had been cached for the survey-prevalence check of §4, and the readouts are the body’s, refitted on the new training split.

Two properties still differ between the classes besides rotation, both astrophysical rather than sampling artefacts: 10.4% of the rotators and 0.5% of the non-rotators carry a flare label, and 0.3% of the rotators against 9.4% of the non-rotators are oscillating giants or solar-like oscillators. The headline rows keep every star, because a survey contains them; two sensitivity rows refit the readout with flare stars removed from both classes, and with the positives restricted to stars in no non-rotation catalogue and the negatives to stars in no catalogue at all.

Result. Table 6 gives both tasks on the first pool and on the rotation pool, every column of Table 1 included. Three things change and one does not. The engineered-feature score on the rotation task falls from 0.540 to 0.386: the first pool’s negatives are stars in no catalogue, the rotation pool’s negatives are every non-rotator in the survey, and the second is the harder set. The fusion gain on rotation rises from $+0.016 \pm 0.016$ to $+0.055 \pm 0.006$, but an untrained encoder now also gains $+0.029$ (init spread $+0.027 \pm 0.003$ over six inits) where on the first pool it lost 0.026; with 74,398 training stars, 128 random columns are no longer redundant to a linear fit. The trained encoder is ahead of the untrained one by $+0.029 \pm 0.008$, paired by seed. Because the pre-registered rule for these two rows requires the untrained control to stay within 2 SE of zero, rotation does not count,

as it did not before. Rotation period does count on both pools, and more clearly on the new one: $+0.034 \pm 0.002$ on 1,318 test stars against $+0.013 \pm 0.010$ on 150, with the untrained control at -0.003 . The body’s count of 7 of 10 is therefore unchanged. The two sensitivity rows agree with the headline on both verdicts. The dynamics-off contrast keeps its sign on both tasks ($+0.019 \pm 0.007$ and $+0.020 \pm 0.003$), and a paired star-bootstrap of every contrast in the table (1,000 resamples, as in Appendix G) agrees with every seed-band verdict.

Two of the body’s other tallies would change if the rotation pool replaced the first pool in Table 1, and are stated here so the reader can recount. Under XGBoost the fusion gain on rotation period turns from $+0.036$ to -0.004 ± 0.003 (the features alone reach 0.829), so Appendix D’s tally of 4 wins, 5 ties and 1 loss becomes 3, 5 and 2. Against the raw-flux MLP the fusion probe is behind on rotation period (0.752 against 0.773, -0.021 ± 0.004), so “ahead on 9 of 10” becomes 8 of 10; it stays ahead of the convolutional baseline on both tasks. Under the linear readout the untrained control’s largest gain across the ten tasks becomes the $+0.029$ above rather than the $+0.016$ the body quotes; that sentence’s conclusion, that the gains which count are not what 128 arbitrary columns buy, is what the rule of this appendix enforces explicitly.

Table 6: The two rotation tasks on the first pool (the body’s numbers) and on the rotation pool. Linear and XGBoost readouts and the two supervised baselines are those of Table 1; gain = features $\oplus \mu$ minus features, mean over six encoder seeds, bold beyond 2 SE; untrained gain = the same with the untrained encoder. Sensitivity rows refit the linear readout on the restricted population. n : test stars; prev.: rotator prevalence.

task	population	n (prev.)	linear readout					XGBoost		supervised	
			features	μ	fusion	gain	untr. gain	features	fusion	Conv1D	MLP
rotation (PR-AUC)	first pool	2021 (0.089)	0.540	0.559	0.555	$+0.016$	-0.026	0.596	0.576	0.565	0.526
	rotation pool	15944 (0.130)	0.386	0.373	0.441	$+0.055$	$+0.029$	0.518	0.515	0.409	0.420
	no flare stars	15670 (0.119)	0.356	0.352	0.413	$+0.057$	$+0.029$	–	–	–	–
	strict	14182 (0.129)	0.367	0.372	0.427	$+0.061$	$+0.036$	–	–	–	–
rotation period (R^2)	first pool	150	0.703	0.677	0.717	$+0.013$	-0.057	0.695	0.730	0.597	0.698
	rotation pool	1318	0.718	0.636	0.752	$+0.034$	-0.003	0.829	0.825	0.697	0.773
	no flare stars	1174	0.713	0.628	0.744	$+0.031$	-0.006	–	–	–	–
	strict	1145	0.718	0.636	0.748	$+0.031$	-0.004	–	–	–	–